

Deep RL

Deep RL = RL + neural networks as function approximators

The methods described in this chapter
find an optimal policy for an
unknown, continuous MDP

- Unknown: states/actions/transitions are not known a-priori
→ agent must rely on trial-and-error (run episodes and see how well it goes)
- Continuous: Too many states/actions to list them in tabular form

Four methods are presented

- Deep Sarsa, Deep Q-learning
- DQN
- Actor-critic

Different types of unknown, discrete/continuous MDPs
(with corresponding algorithm and example)

	discrete action space	continuous action space
discrete state space	SARSA / Q-learning / MC (robot maze)	actor-critic
continuous state space	Deep Sarsa, Deep Q-Learning, DQN (Breakout, Space Invaders)	actor-critic (autonomous driving)

Motivation/explanation 1 for Deep Sarsa, Deep Q-learning

- Recap: SARSA update rule

$$Q(s, a) \leftarrow Q(s, a) + \alpha(R + \gamma Q(s', a') - Q(s, a))$$

- Recap: Q-learning update rule

$$Q(s, a) \leftarrow Q(s, a) + \alpha \left(R + \gamma \max_{a'} Q(s', a') - Q(s, a) \right)$$

- What happens after we calculated the optimal Q-function?

→ There is a unique optimal Q-function, so it will not change

→ SARSA / Q-learning update rules will be of the form

$$Q(s, a) \leftarrow Q(s, a)$$

→ the other terms in the SARSA/Q-learning update rule must be zero, e.g. it is for SARSA

$$R + \gamma Q(s', a') - Q(s, a) = 0$$

Motivation/explanation 2 for Deep Sarsa, Deep Q-learning

- Approximate the true Q-function $Q(s, a)$ with a neural network $\hat{Q}(s, a, \mathbf{w})$
- Minimize loss L measuring the difference between true and approximated Q-function through gradient descent

- Define loss function

$$L = \left(Q(s, a) - \hat{Q}(s, a, \mathbf{w}) \right)^2$$

- Calculate gradient

$$\nabla_{\mathbf{w}} L = -2 \left(Q(s, a) - \hat{Q}(s, a, \mathbf{w}) \right) \nabla_{\mathbf{w}} \hat{Q}(s, a, \mathbf{w})$$

- Use approximation of the true Q-function, e.g. for SARSA

$$\nabla_{\mathbf{w}} L = -2 \left(R + \gamma \hat{Q}(s', a', \mathbf{w}) - \hat{Q}(s, a, \mathbf{w}) \right) \nabla_{\mathbf{w}} \hat{Q}(s, a, \mathbf{w})$$

- Update weights \mathbf{w} through (stochastic) gradient descent

Deep Sarsa

- Train a neural network to approximate the true Q-function $Q(s, a)$ with $\hat{Q}(s, a, \mathbf{w})$
- Loss L to be minimized is of the form

$$L = \left(R + \gamma \hat{Q}(s', a', \mathbf{w}) - \hat{Q}(s, a, \mathbf{w}) \right)^2$$

- If the loss becomes zero for all states/actions, the optimal Q-function is obtained (Q-function won't be updated anymore)

~~$$\hat{Q}(s, a, \mathbf{w}) \leftarrow \hat{Q}(s, a, \mathbf{w}) + \alpha \left(R + \gamma \hat{Q}(s', a', \mathbf{w}) - \hat{Q}(s, a, \mathbf{w}) \right)$$~~

- Resulting gradient becomes

$$\nabla_{\mathbf{w}} L = -2 \left(R + \gamma \hat{Q}(s', a', \mathbf{w}) - \hat{Q}(s, a, \mathbf{w}) \right) \nabla_{\mathbf{w}} \hat{Q}(s, a, \mathbf{w})$$

Deep Q-learning

- Train a neural network to approximate the true Q-function $Q(s, a)$ with $\hat{Q}(s, a, \mathbf{w})$
- Loss L to be minimized is of the form

$$L = \left(R + \gamma \max_{a'} \hat{Q}(s', a', \mathbf{w}) - \hat{Q}(s, a, \mathbf{w}) \right)^2$$

- If the loss becomes zero for all states/actions, the optimal Q-function is obtained (Q-function won't be updated anymore)


$$\hat{Q}(s, a, \mathbf{w}) \leftarrow \hat{Q}(s, a, \mathbf{w}) + \alpha \left(R + \gamma \max_{a'} \hat{Q}(s', a', \mathbf{w}) - \hat{Q}(s, a, \mathbf{w}) \right)$$

- Resulting gradient becomes

$$\nabla_{\mathbf{w}} L = -2 \left(R + \gamma \max_{a'} \hat{Q}(s', a', \mathbf{w}) - \hat{Q}(s, a, \mathbf{w}) \right) \nabla_{\mathbf{w}} \hat{Q}(s, a, \mathbf{w})$$

How to train Deep Sarsa / Deep Q-learning

1. Run certain number of episodes with neural network (type shown below)
2. Store data of the form (s, a, r, s', a') for every time step in a buffer
3. Train neural network to minimize the loss L , then clear the buffer and go to 1.
(Neural network has one output per action, representing its probability)



Note that finding a ϵ -greedy/optimal policy based on the obtained values $\hat{Q}(s, a_1, w) \dots \hat{Q}(s, a_n, w)$ is easy, just pick the action with the largest value

DQN (Deep Q Network)

- Deep Q-learning with improvements:
 - experience replay
 - target network
- first "proper" deep RL approach
(deep SARSA and deep Q-learning are just for didactic purpose)

Replay	○	○	×	×
Target	○	×	○	×
Breakout	316.8	240.7	10.2	3.2
River Raid	7446.6	4102.8	2867.7	1453.0
Seaquest	2894.4	822.6	1003.0	275.8
Space Invaders	1088.9	826.3	373.2	302.0

<https://towardsdatascience.com/welcome-to-deep-reinforcement-learning-part-1-dqn-c3cab4d41b6b>

Experience replay

Idea: Use all data obtained so far for training

Training steps

1. Run certain number of episodes with neural network
2. Store data of the form (s, a, r, s', a') for every time step in a buffer
3. Train neural network to minimize the loss L , ~~then clear the buffer~~ and go to 1.

Implementation

- buffer often has finite size, delete oldest data if buffer is full
- training often happens only on a (random) subset of all buffer data (faster if buffer is large)

Effect of experience replay:

- Is possible, because Q-learning is off-policy
- Each observed transition (s, a, r, s', a') can be used multiple times for training
→ better data efficiency
- History contains data from many different policies
→ lots of different samples seen

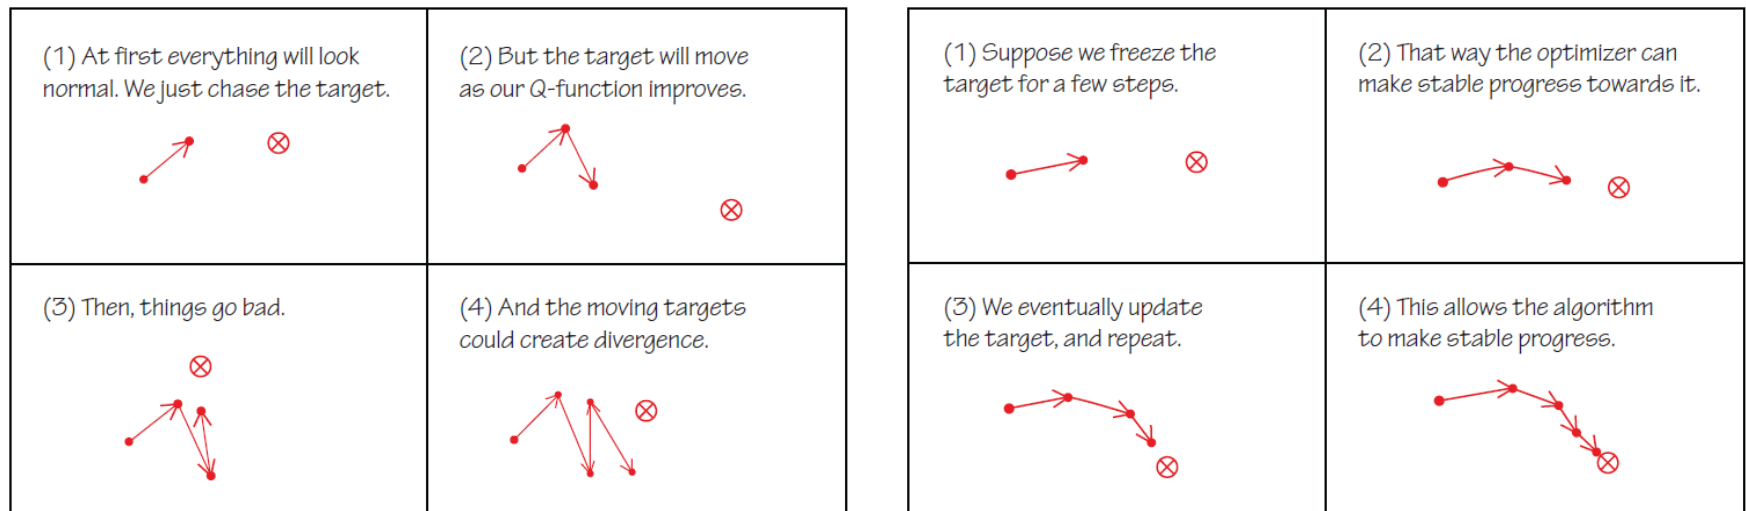


Task: What does off-policy mean?

Target network

Idea: Fix the **orange term (target network)** for a certain number of training steps to avoid unstable behaviour

$$\hat{Q}(s, a, \mathbf{w}) \leftarrow \hat{Q}(s, a, \mathbf{w}) + \alpha \left(R + \gamma \max_{a'} \hat{Q}(s', a', \mathbf{w}) - \hat{Q}(s, a, \mathbf{w}) \right)$$



<https://livebook.manning.com/concept/reinforcement-learning/target-network>

Effect of target networks:

- By using target networks, the goal $R + \gamma \max_{a'} \hat{Q}(s', a', \mathbf{w})$ that should be learned by $\hat{Q}(s, a, \mathbf{w})$ in order to minimize

$$L = \left(R + \gamma \max_{a'} \hat{Q}(s', a', \mathbf{w}') - \hat{Q}(s, a, \mathbf{w}) \right)^2$$

changes only slowly \rightarrow more stable learning

Training steps with experience replay and target network

1. Initialization: Create two identical neural networks $\hat{Q}(s, a, \mathbf{w})$ and $\hat{Q}(s, a, \mathbf{w}')$
2. Repeat for a certain number of times
 - a) Run certain number of episodes (policy based on neural network $\hat{Q}(s, a, \mathbf{w})$)
 - b) Store data of the form (s, a, r, s', a') for every time step in a buffer
 - c) Train neural network $\hat{Q}(s, a, \mathbf{w})$ to minimize the loss L

$$L = \left(R + \gamma \max_{a'} \hat{Q}(s', a', \mathbf{w}') - \hat{Q}(s, a, \mathbf{w}) \right)^2$$

This corresponds to the update rule


$$\hat{Q}(s, a, \mathbf{w}) \leftarrow \hat{Q}(s, a, \mathbf{w}) + \alpha \left(R + \gamma \max_{a'} \hat{Q}(s', a', \mathbf{w}') - \hat{Q}(s, a, \mathbf{w}) \right)$$

3. Copy $\hat{Q}(s, a, \mathbf{w})$ to obtain $\hat{Q}(s, a, \mathbf{w}')$ and go to 2.

Task: Why is there no Deep SARSA with experience replay and target networks?

Example: Playing Atari with Deep RL (2013)

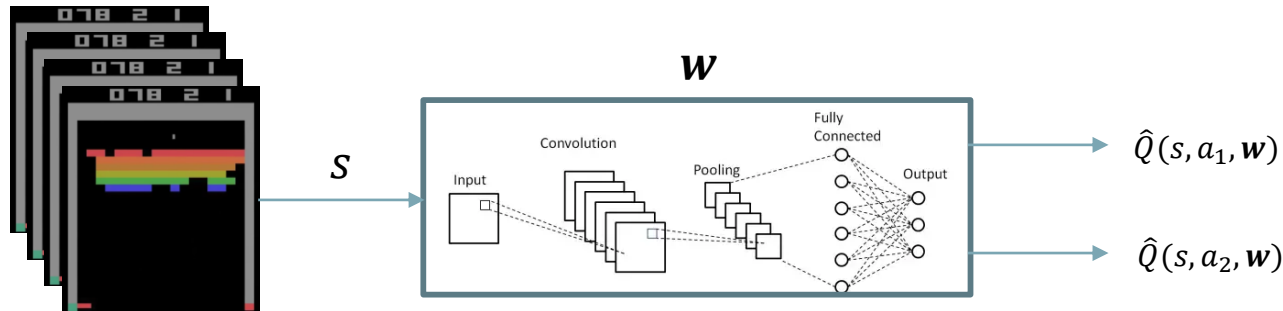
- "inventor" of DQN
- breakthrough as it uses raw images as input
- no domain knowledge
→ generalizes to other tasks
- 2015: presented in Nature
- paved the road for many modern deep RL techniques
- successor: DDQN



<https://www.youtube.com/watch?v=TmPTTjtdgg>

Deep RL

- function approximator is a convolutional neural network
- state is encoded by the last four images



Task: What do the two outputs $\hat{Q}(s, a_1, \mathbf{w})$ and $\hat{Q}(s, a_2, \mathbf{w})$ represent?

Task: Why are the last four images used to represent the state (and not just a single one)?

DDQN (Double Deep Q Networks)

- Problem of DQN (similar to Q-Learning): max-operator in

$$\hat{Q}(s, a, \mathbf{w}) \leftarrow \hat{Q}(s, a, \mathbf{w}) + \alpha \left(R + \gamma \max_{a'} \hat{Q}(s', a', \mathbf{w}') - \hat{Q}(s, a, \mathbf{w}) \right)$$

may overestimate the true Q-value if \hat{Q} is inaccurate

- There are two situations where this has an influence
 - selecting an action during episodes
 - rating an action during training
- Idea of DDQN: Decouple these two steps in to reduce the effect of a wrong estimate (use two different networks)

- Original update rule (target)

$$\hat{Q}(s, a, \mathbf{w}) \leftarrow \hat{Q}(s, a, \mathbf{w}) + \alpha \left(R + \gamma \max_{a'} \hat{Q}(s', a', \mathbf{w}') - \hat{Q}(s, a, \mathbf{w}) \right)$$

- DQN has the target

$$R + \gamma \max_{a'} \hat{Q}(s', a', \mathbf{w}')$$

$$R + \gamma \hat{Q}(s', \operatorname{argmax}_{a'} \hat{Q}(s', a', \mathbf{w}'), \mathbf{w}')$$



- DDQN has the target

$$R + \gamma \hat{Q}(s', \operatorname{argmax}_{a'} \hat{Q}(s', a', \mathbf{w}'), \mathbf{w})$$

Dueling (D)DQN

- Problem of DQN/DDQN: The absolute Q-values are not that important when choosing an action, it is more their relative value w.r.t. each other
- Idea: Split up the learned Q-function into two different networks:
 - One network can learn the absolute value
 - The other network can learn the relative value for different actions
- Formula used in [paper](#)

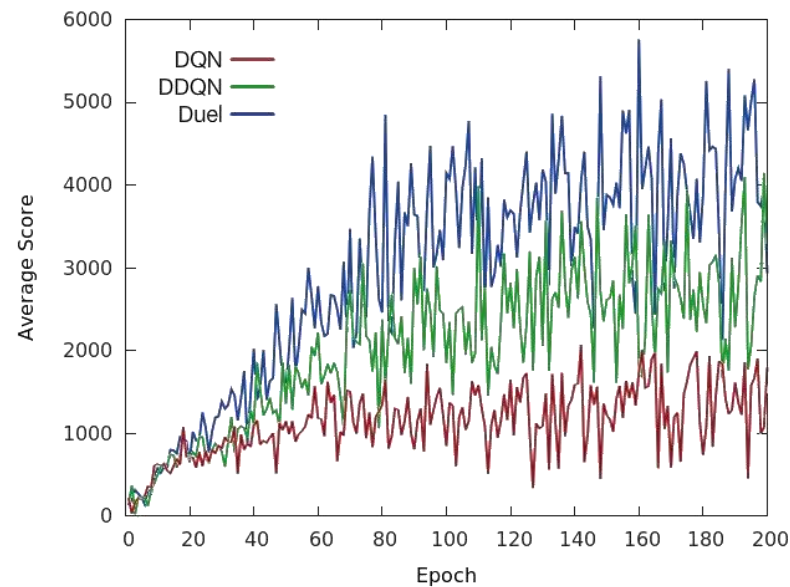
$$\hat{Q}(s, a, \mathbf{w}) = V(s, \mathbf{w}) + \left(A(s, a, \mathbf{w}') - \max_a A(s, a, \mathbf{w}') \right)$$



Deep RL

Example: Space Invaders


- performance of DQN / DDQN / Dueling DQN over trained epoch



http://torch.ch/blog/2016/04/30/dueling_dqn.html

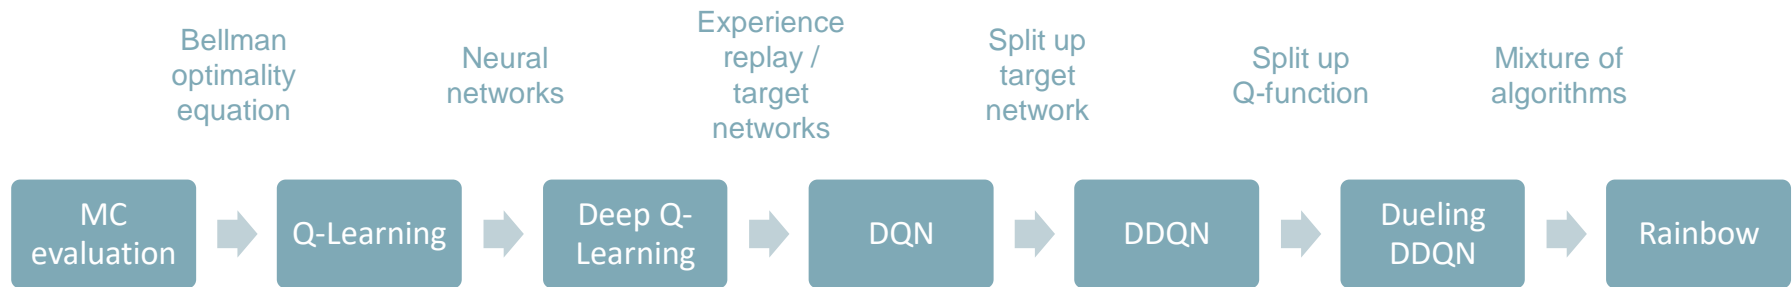
Rainbow

- Idea: Because different DQN variants excel for certain problems, iterate over all of them during training



<https://www.arxiv-vanity.com/papers/1710.02298/>


Evolution of RL algorithms



Task: What are the improvements for each RL algorithm?

Actor-critic methods

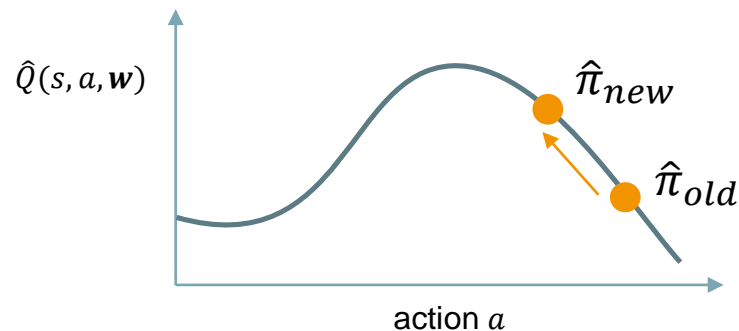
- are used mainly for continuous state / continuous action MDPs (\leftrightarrow DQN: continuous state / discrete action)
- rely on two neural networks
 - critic: approximate optimal (V-)/Q-function based on current state/action
 - actor: approximate optimal policy based on current state



- the critic $\hat{Q}(s, a, \mathbf{w})$ is trained to learn the optimal (V-)/Q-function, e.g. by minimizing the loss L


$$L = \left(R + \gamma \hat{Q}(s', a', \mathbf{w}) - \hat{Q}(s, a, \mathbf{w}) \right)^2$$

- the actor $\hat{\pi}(s, \mathbf{w}')$ is trained through gradient ascent on the critic to output the optimal action for a given state such that the critic output $\hat{Q}(s, \hat{\pi}(s, \mathbf{w}'), \mathbf{w})$ is maximized



State-of-the-art RL algorithms


- more additional improvements (similar to experience replay / target networks) lead to state-of-the-art algorithms (e.g. A2C, A3C, DDPG, PPO, TRPO, TD3, SAC, ...)
- rule of thumb: The newer the algorithm is, the better it performs (learns faster)
- Overview of different algorithms:
 - <https://lilianweng.github.io/posts/2018-04-08-policy-gradient/>
 - <https://medium.datadriveninvestor.com/which-reinforcement-learning-rl-algorithm-to-use-where-when-and-in-what-scenario-e3e7617fb0b1>
 - https://www.ias.informatik.tu-darmstadt.de/uploads/Team/DavideTateo/felix_thesis.pdf



<https://www.arxiv-vanity.com/papers/1909.01500/>

Brief summary

- Deep RL is used whenever the state/action space is continuous
- Deep Sarsa, Deep Q-learning are not used in practice
- DQN is used for continuous state / discrete action spaces
- DQN approximates the Q-function with a neural network
- Actor-critic methods are used for continuous state / continuous action spaces
- Actor-critic methods approximate the Q-function and the policy with neural networks



<https://www.freecodecamp.org/news/an-intro-to-advantage-actor-critic-methods-lets-play-sonic-the-hedgehog-86d6240171d/>

Kahoot!

Kahoot!