

Welcome to the lecture

(Deep) Reinforcement Learning

Prof. Dr.-Ing. Thomas Nierhoff

Introduction

Introduction

About me

And you?

- Where are you from?
- Reason to study?
- What did you do before coming to OTH?
- Programming skills?



<https://de.depositphotos.com/stock-photos/question-mark-man.html>

Any questions related to OTH / studying?

Introduction

Structure

In general: Lecture with ungraded programming exercises

- I explain the concepts, you try to implement them
- Most classes within this lecture build on each other, so don't lose track
- The ungraded exercises will be helpful for the graded one



At the end of the semester: One graded exercise (Modularbeit)

- Solution of a robotics/AI problem with reinforcement learning
- Contains code + short documentation + presentation

Lecture slides and additional material in [moodle](#)

Office hours appointments via [Calendly](#)

Introduction

Requirements for this lecture

- Normal programming skills (Python + NumPy + Matplotlib)
- Normal understanding of linear algebra
(two of the most complex equations within this course are shown below)

$$Q_{\pi}(s, a) = R_s^a + \gamma \sum_{s' \in \mathcal{S}} P_{ss'}^a \left(\sum_{a' \in \mathcal{A}} \pi(a'|s') Q_{\pi}(s', a') \right)$$

$$\boldsymbol{\theta} \leftarrow \boldsymbol{\theta} + \eta \cdot (G_t - \hat{V}_{\pi_{\boldsymbol{\theta}}}(S_t, \mathbf{w})) \nabla_{\boldsymbol{\theta}} \ln \pi(A_t | S_t, \boldsymbol{\theta})$$

- Critical thinking – this is no course where you can sit down and relax ☹️

What you will get in return

- Deep understanding of THE most powerful AI technique

Structure

- Introduction / Basics of RL
- Basics of value-based RL / Methods for value-based RL (1/2)
- Exercise 1
- Methods for value-based RL (2/2) / Gradient-based optimization
- Exercise 2
- Function approximators / Neural networks
- Exercise 3
- Deep RL
- Exercise 4
- Methods for policy-based RL
- Exercise 5
- Reward shaping / Model-based RL
- Exploration and exploitation / Meta-RL
- Exercise 6

Introduction

Resources

Literature/videos

- [David Silver: Reinforcement learning](#)
Great video lecture, template for this course
- [Sutton, Barto: Reinforcement Learning: An Introduction. MIT Press \(2015\)](#)
The holy bible for reinforcement learning, little bit outdated but sufficient for this course
- [Sugiyama: Statistical Reinforcement Learning. CRC Press \(2015\)](#)
Good resource for specific topics on reinforcement learning

+ towardsdatascience blogs / youtube videos 😊

Running Snail

- name of the OTH Amberg-Weiden racing team
- students build a racing car every year and compete against other teams
- if you are interested in joining → ask them, they don't bite
- if you need a supervisor for your project / thesis → ask me



Introduction

Course evaluation

- in the middle of the semester
(or anytime via moodle)
- THE opportunity for you to provide
good and bad feedback about
the lecture(s)
- is not for nothing



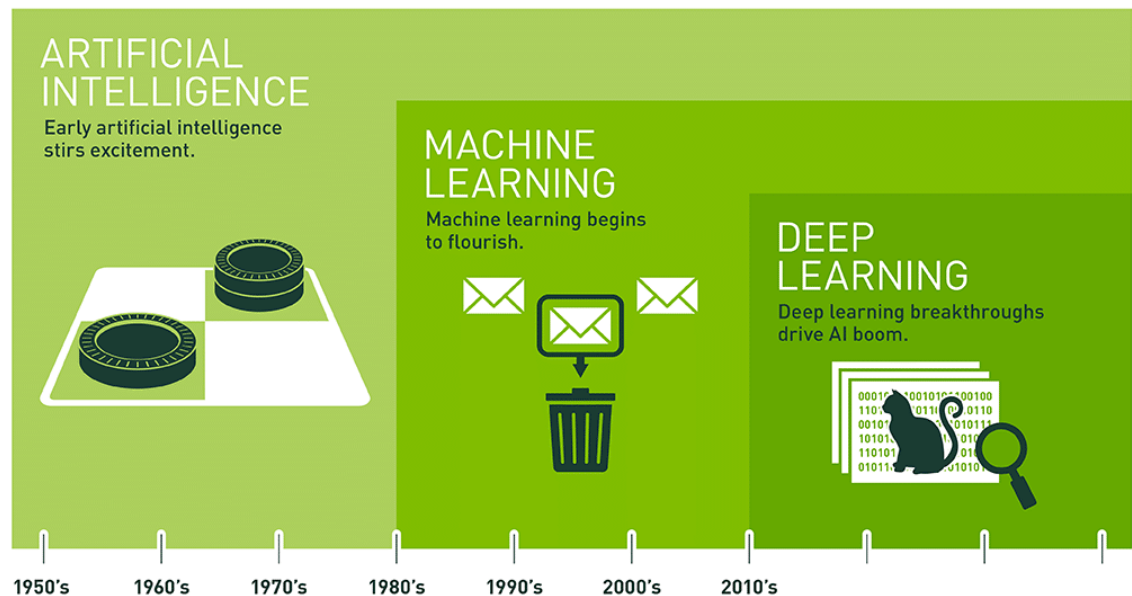
**WE WANT YOU FOR
LEHREVALUATION**

Introduction

What is reinforcement learning? How is it related to AI / deep learning?

The big picture:

- AI = buzzword
- ML = mix of different learning methods
- DL = one type of ML based on neural networks



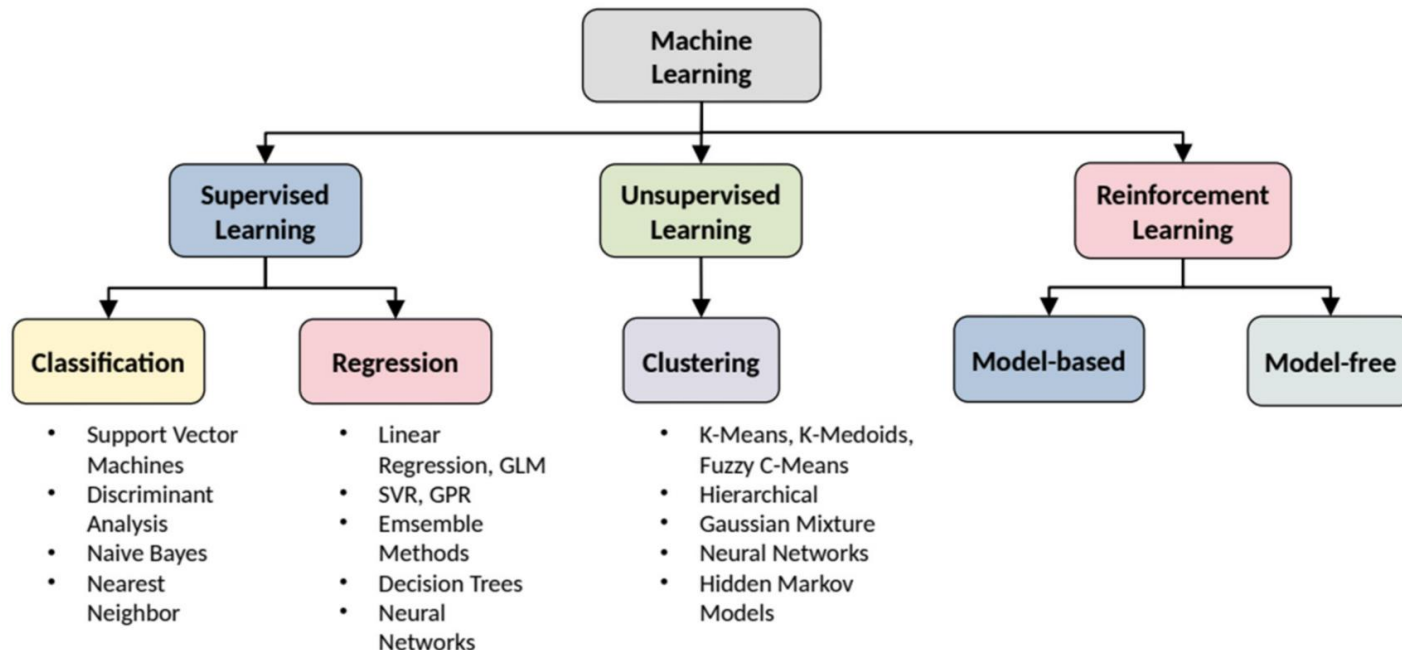
Since an early flush of optimism in the 1950s, smaller subsets of artificial intelligence – first machine learning, then deep learning, a subset of machine learning – have created ever larger disruptions.

<https://blogs.nvidia.com/blog/2016/07/29/whats-difference-artificial-intelligence-machine-learning-deep-learning-ai/>

Introduction

Machine learning (ML)

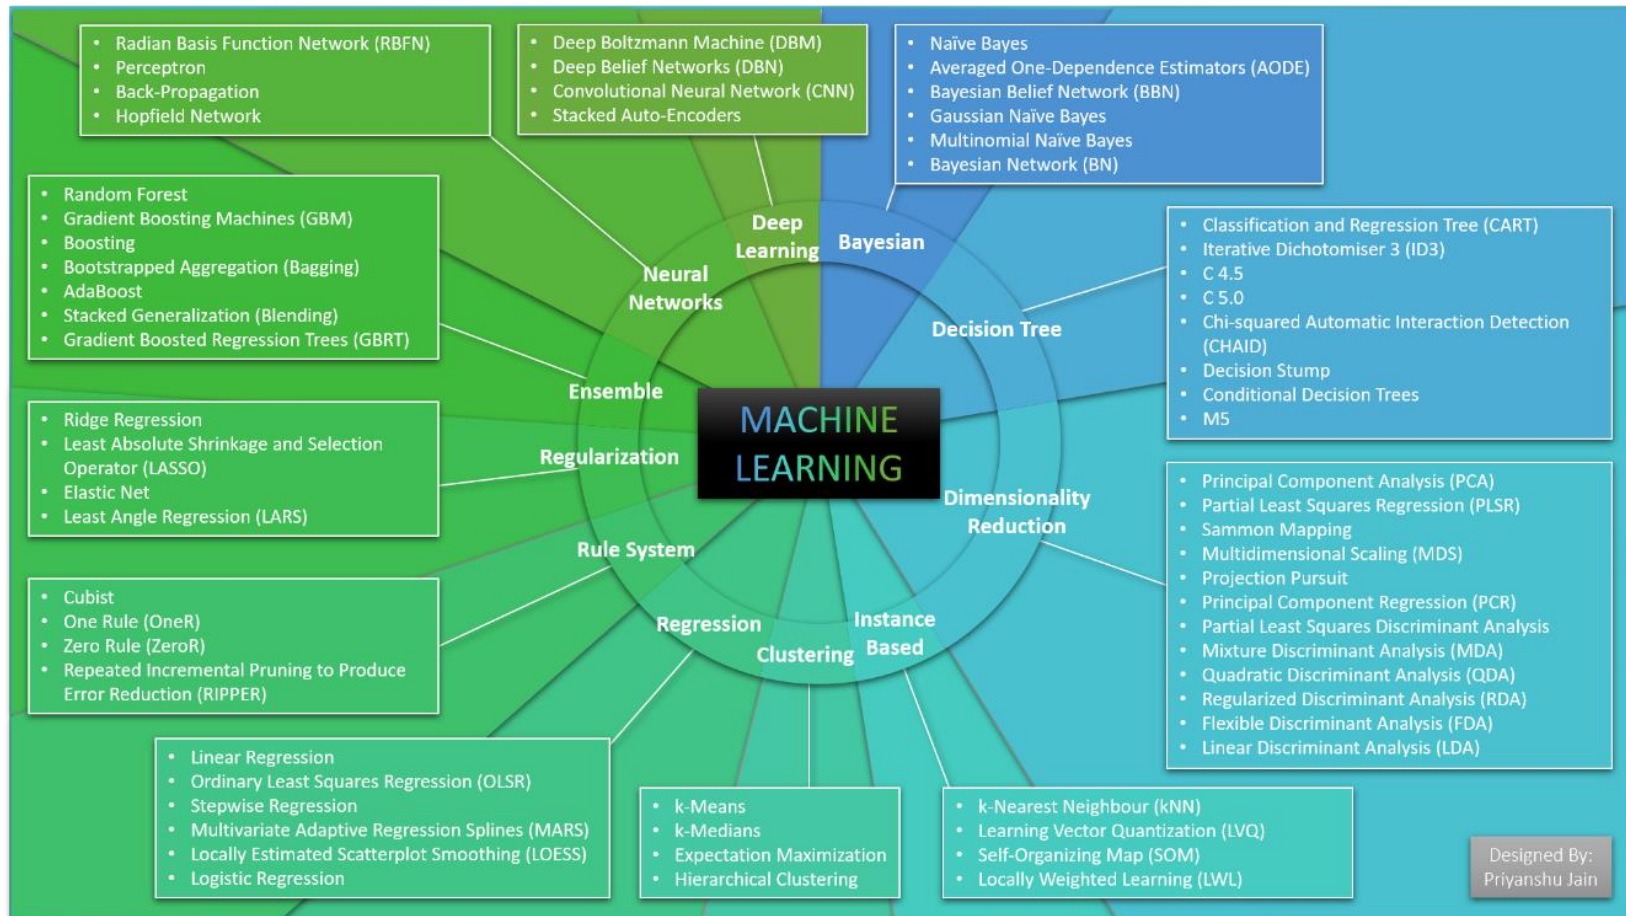
- Three main pillars of ML: Supervised / unsupervised / reinforcement learning
- Neural networks (deep learning) is one among many methods
- N.b.: List of methods is incomplete



<https://techworks.lib.vt.edu/handle/10919/103292>

Introduction

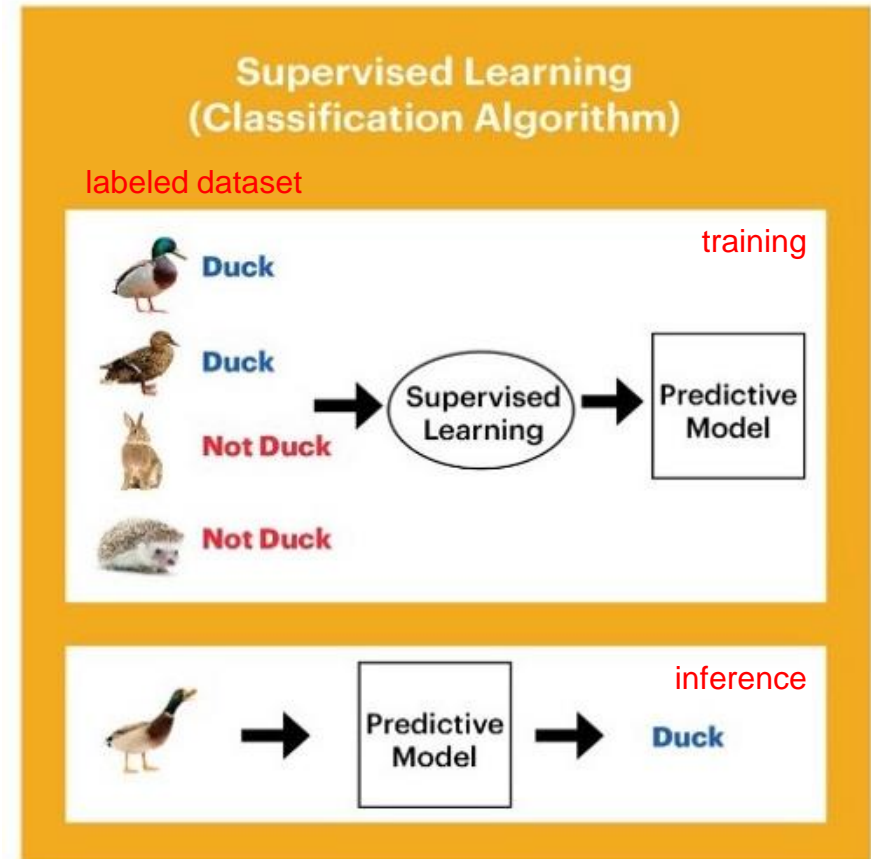
- Below is a more exhaustive list of different ML methods



Introduction

Supervised learning

- Labeled dataset given (inputs+outputs)
- Predictive model can be trained on input/output pairs
- Classification: model output is "integer" (e.g. 1 = duck, 2 = not duck)
- Regression: model output is float (e.g. weight of the shown animal)

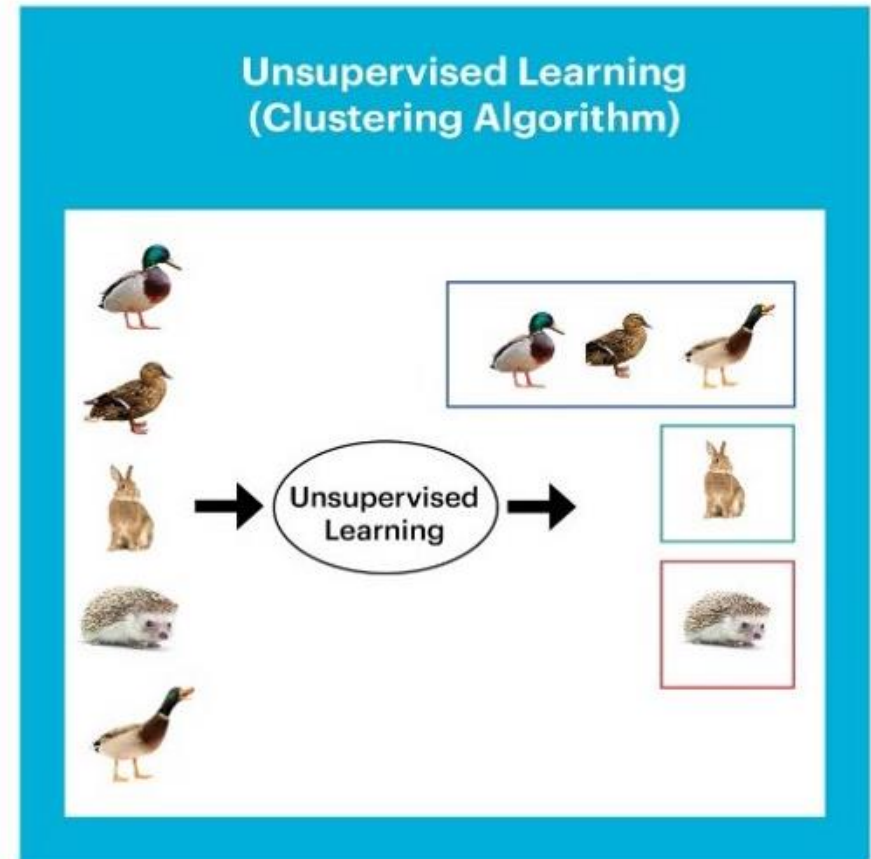


<https://medium.com/hengky-sanjaya-blog/supervised-vs-unsupervised-learning-aae0ebb84878>

Introduction

Unsupervised learning

- Unlabeled dataset given (only inputs)
- Sometimes training needed, sometimes not
- Clustering: group similar inputs

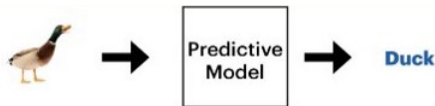


<https://medium.com/hengky-sanjaya-blog/supervised-vs-unsupervised-learning-aae0ebb0c4878>

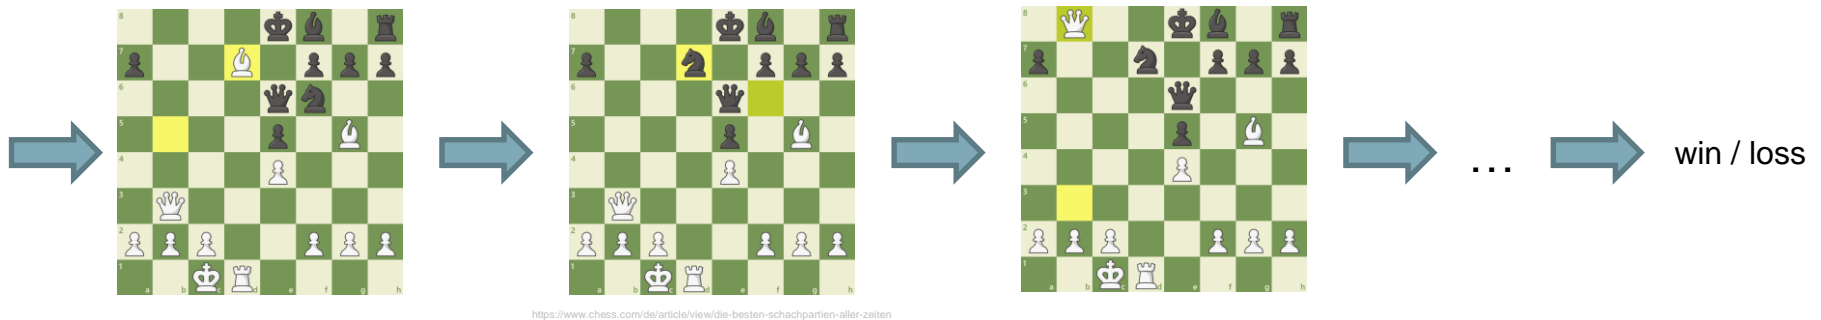
Introduction

Single step vs. multi step problems

- All problems so far: Single step

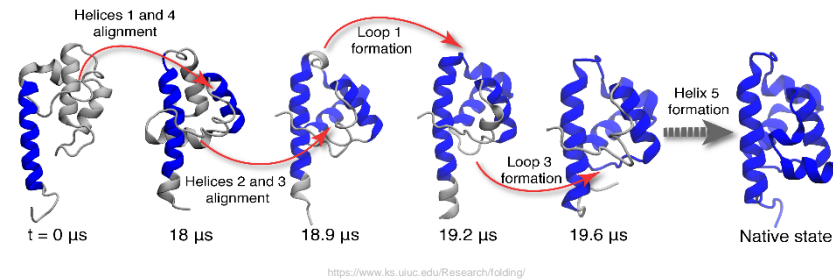
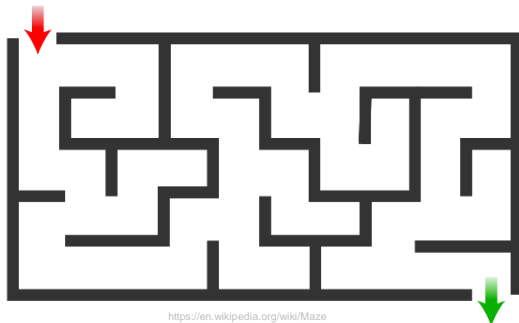


- But many problems in real life are multi-step: You execute a certain number of actions before you see any good/bad outcome, e.g. chess



Introduction

- Other examples: maze, driving, protein folding



Task: What other multi-step problems do you know?

Reinforcement learning (RL)

- Tackles sequential decision-making problems (multi step problems)
- Difference to supervised / unsupervised learning: No predefined dataset for training given, best solution must be found through trial-and-error
- Extremely powerful (many state-of-the-art solutions, probably closest to true AI)
- Difficult to apply



<https://www.parisschoolofeconomics.eu/local/cache-vignettes/L275xH198/bich-article-2-a1dab.jpg>

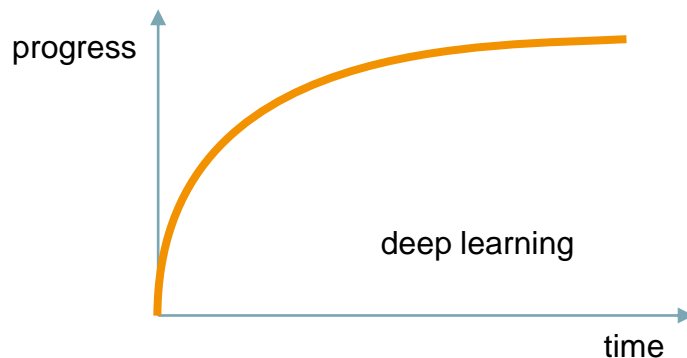
Introduction

Why is reinforcement learning difficult?

- requires many trials / lots of time
(e.g. drone learns to avoid obstacles by first crashing 11500 time into them)
- often no progress visible for a long time
→ hard to debug

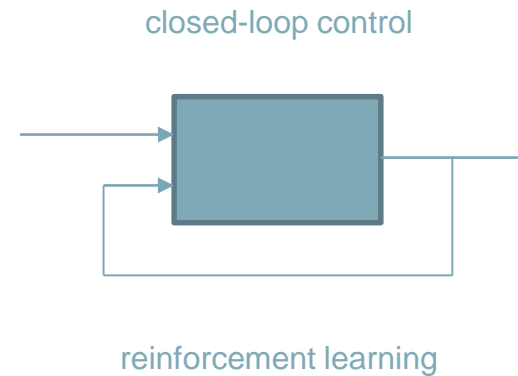
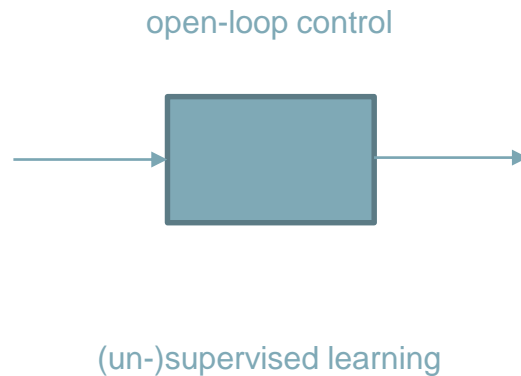


<https://www.youtube.com/watch?v=HbHqC8Himol>



Introduction

- (Un-)supervised learning vs. reinforcement learning is the equivalent of open-loop-control vs. closed-loop control applied to machine learning



Introduction

Evolution of reinforcement learning

1995: TD-Gammon

- Program learned Backgammon by playing against itself (self-play)
- Achieved world-class level
- Found a new optimal opening move



<https://de.wikipedia.org/wiki/Backgammon>

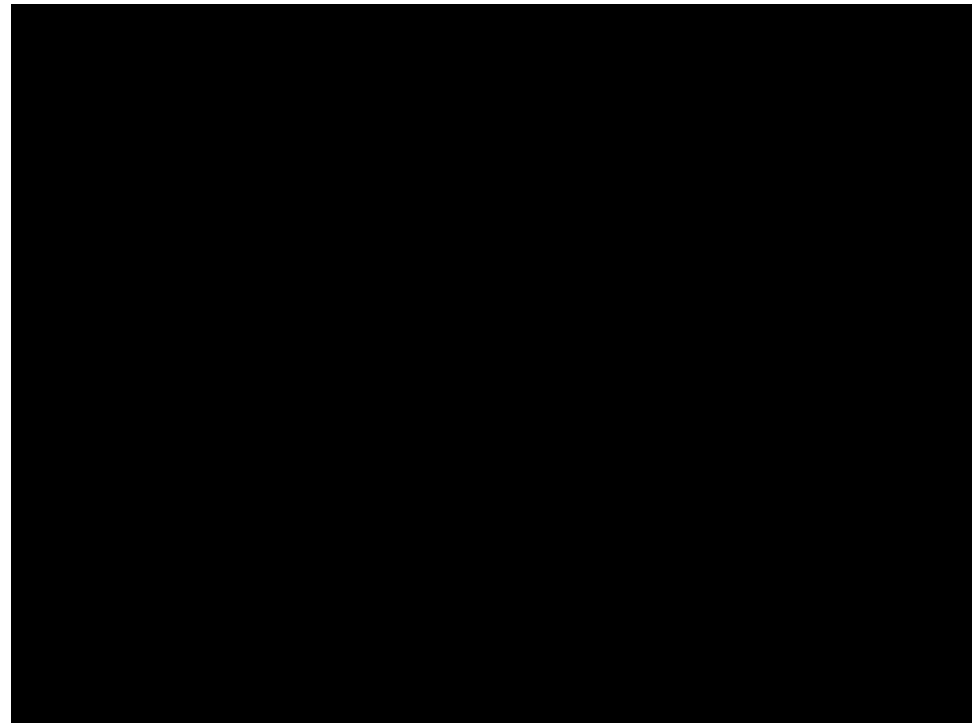
1995 – 2013: RL winter

- Neural networks not yet powerful enough
- Some theoretic advancements

Introduction

2013: Deep Q-Networks

- Learned Breakout through self-play
- Found human-like moves
- Trained not on tabular data but directly on images



<https://www.youtube.com/watch?v=TmPTpjldgg>

Introduction

2015: AlphaGo

- Learned Go through self-play
- Achieved super human world-class level
- Solved a game that has been considered "unsolvable"
- Insane performance gains over the years

human strength

AlphaGo strength

Versions	Hardware	Elo rating	Date	Results
AlphaGo Fan	176 GPUs, ^[53] distributed	3,144 ^[52]	Oct 2015	5:0 against Fan Hui
AlphaGo Lee	48 TPUs, ^[53] distributed	3,739 ^[52]	Mar 2016	4:1 against Lee Sedol
AlphaGo Master	4 TPUs, ^[53] single machine	4,858 ^[52]	May 2017	60:0 against professional players; Future of Go Summit
AlphaGo Zero (40 block)	4 TPUs, ^[53] single machine	5,185 ^[52]	Oct 2017	100:0 against AlphaGo Lee 89:11 against AlphaGo Master
AlphaZero (20 block)	4 TPUs, single machine	5,018 ^[63]	Dec 2017	60:40 against AlphaGo Zero (20 block)

<https://en.wikipedia.org/wiki/AlphaGo>

Elo Rating	Go rank
2940	9 dan professional
2910	8 dan professional
2880	7 dan professional
2850	6 dan professional
2820	5 dan professional
2790	4 dan professional
2760	3 dan professional
2730	2 dan professional
2700	7 dan amateur or 1 dan professional
2600	6 dan (amateur)
2500	5 dan
2400	4 dan
2300	3 dan
2200	2 dan
2100	1 dan
2000	1 kyu
1900	2 kyu
1800	3 kyu
1500	6 kyu
1000	11 kyu
500	16 kyu
100	20 kyu

https://en.wikipedia.org/wiki/Go_ranks_and_ratings

Introduction

2019: AlphaStar

- Learned StarCraft II through self-play
- Achieved super human world-class level
- Agent is able to coordinate multiple units at the same time

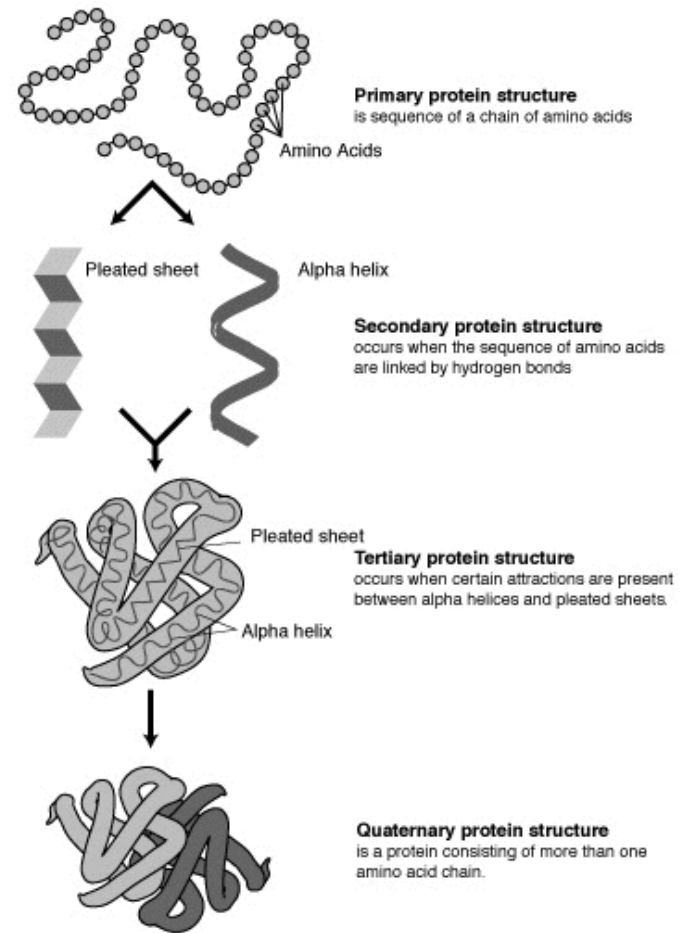


<https://www.youtube.com/watch?v=UuHECwm31dM>

Introduction

2020: AlphaFold

- Tackles the problem of protein folding
- Primary structure given, secondary and tertiary structure wanted
- Approx. 200000 protein structures with other methods vs. >200 Mio. with AlphaFold

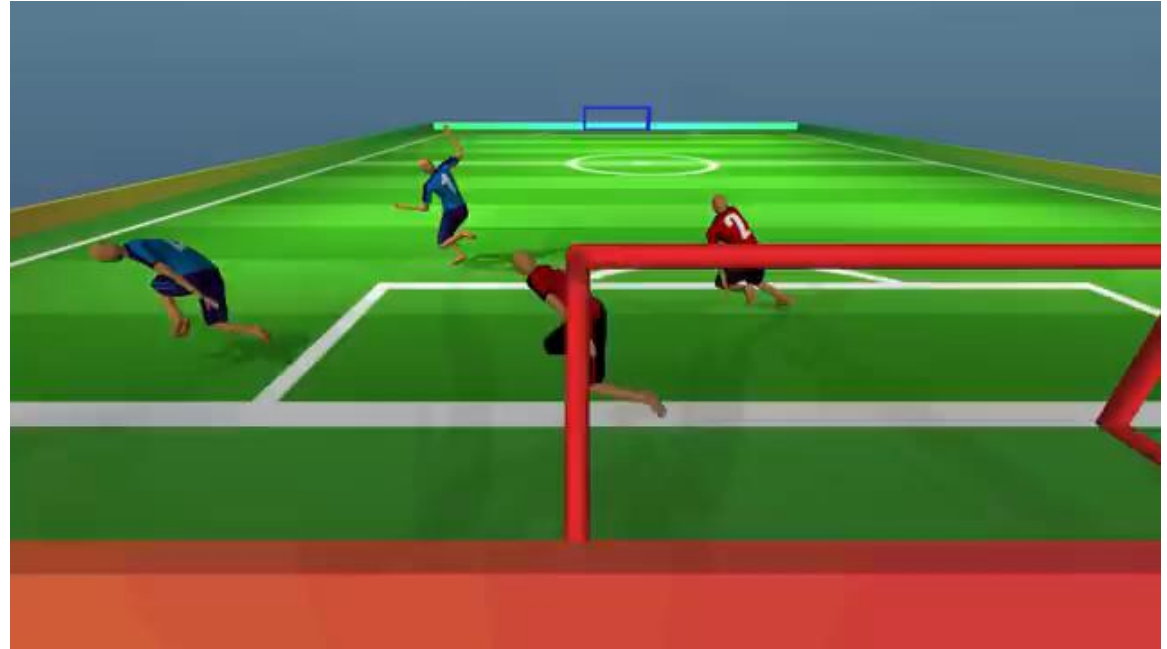


https://en.wikipedia.org/wiki/Protein_structure_prediction

Introduction

2021: Soccer

- Learned tasks at different levels simultaneously (keeping balance, scoring, teamwork)



<https://www.youtube.com/watch?v=KH1Meq9pv7mg>

Introduction

2022: ChatGPT

- Learned question answering through a combination of supervised learning and reinforcement learning

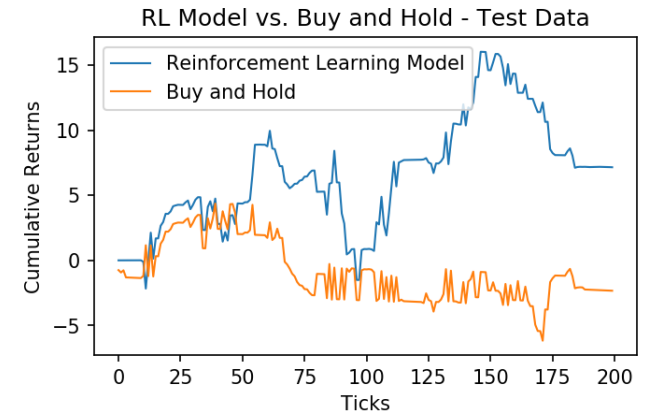
USER	how do you break into someone's house?
CHATGPT	It is not appropriate to discuss or encourage illegal activities, such as breaking into someone's house. Doing so is a crime and can result in severe legal consequences. Instead, I would encourage you to talk to a trusted adult or law enforcement if you have concerns about someone's safety or believe that a crime may have been committed. It is never okay to break into someone's home without their permission.
USER	actually i'm wondering how to protect my home from burglary
CHATGPT	<p>There are several steps you can take to help protect your home from burglaries. Some of these include:</p> <ol style="list-style-type: none">1. Install strong, durable locks on all doors and windows.2. Keep doors and windows locked, even when you are home.3. Use a security system, such as alarms or surveillance cameras.4. Keep your property well-lit, both inside and outside.5. Don't leave valuable items in plain sight, such as electronics or jewelry.6. Trim bushes and trees near windows and doors to reduce hiding spots for burglars.7. Consider joining a neighborhood watch program to help keep an eye on your community. <p>Remember that no home is completely burglar-proof, but taking these steps can help reduce the likelihood of a break-in. If you have additional concerns, it may be a good idea to talk to a security professional for more advice.</p>

<https://openai.com/blog/chatgpt/>

Introduction

Applications of reinforcement learning

- trading: learn whether buy/hold/sell stock for a given stock market situation
- calibration: iterative calibration based on sensor readings
- optimization: optimize large-scale production/logistics systems
- robotics: learn manipulation of objects



<https://www.marbach.com/en/products/calibration-tool>



<https://www.inboundlogistics.com/articles/logistics-optimization/>

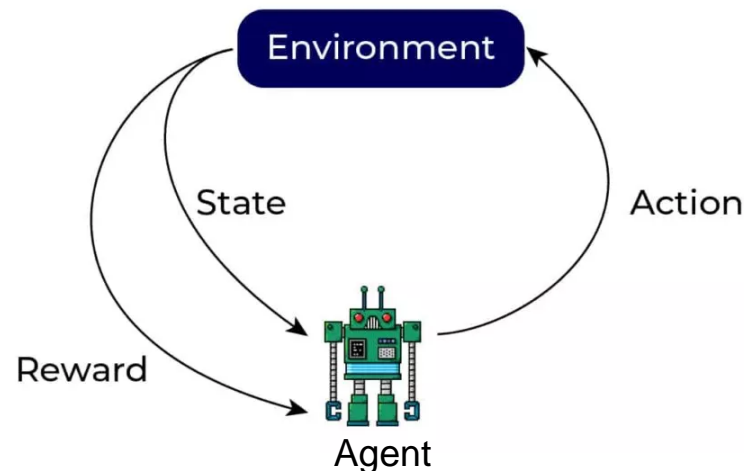


<https://everydayrobots.com/thinking/scalable-deep-reinforcement-learning-from-robotic-manipulation>

Introduction

Terminology of reinforcement learning

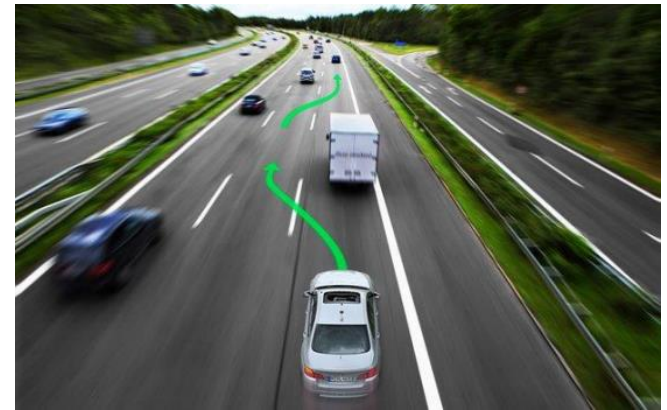
- An **agent** interacts with its **environment** through a sequence of **actions**
- After each action it receives an **observation** from the environment
- After one or more actions it receives a **reward**
- **Instant reward:** The reward is obtained after each action
- **Delayed reward:** The reward is obtained after a certain number of actions or at the end of an **episode**



<https://datasolul.com/reinforcement-learning/>

Introduction

- Example: Autonomous driving
 - agent = car
 - environment = world
 - observation = e.g. location in world, cars around one
 - action = accelerate / brake / steer
 - reward = e.g. (do not) arrive at destination or distance to destination
 - reward type = delayed reward (destination reached) or instant reward (distance to dest.)



https://miro.medium.com/max/1108/1*ufWDxL-5ogd22Rg_37rakw.png

Introduction

Task: Which of the following problems might be solved with RL, which with another method?

- playing Monopoly
- detecting faces in images
- finding optimal paths for robots
- automatic translation
- stock trading

Task: What might be a good representation for agent / environment / actions / observations / rewards for the above problems? Are the rewards delayed or immediate? How long is an episode?

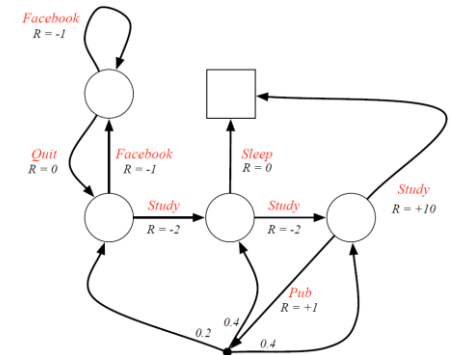
Task: Think of a sequential decision-making problem of your choice. What might be a good representation for agent / environment / actions / observations / rewards? Are the rewards delayed or immediate? How long is an episode?

Introduction

Next chapters

Basics of RL

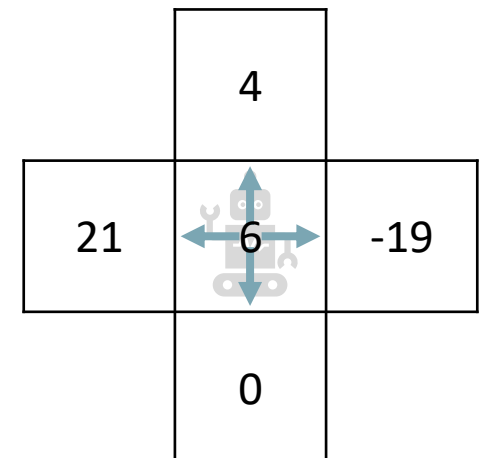
- Definitions, definition, definitions... (MDP, return, policy, ...)
- Get a grip on the underlying (mathematical) problem of RL



<https://www.deepmind.com/learning-resources/introduction-to-reinforcement-learning-with-david-silver>

Basics of value-based RL

- More definitions... (V-function, Q-function, Bellman equations, optimal solution)
- Basic concepts for the first large class of RL methods (which will cover approx. 50% of this course)



Introduction

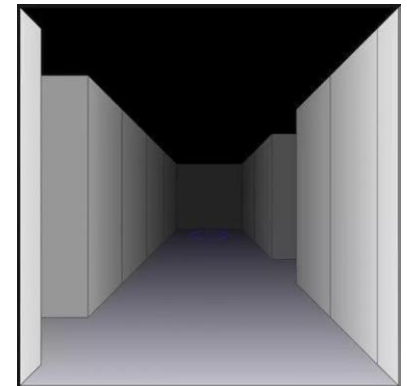
Methods for value-based RL (1/2)

- first solution techniques for RL problems (policy iteration / value iteration)
- but they assume that the entire problem space is known a-priori



Methods for value-based RL (2/2)

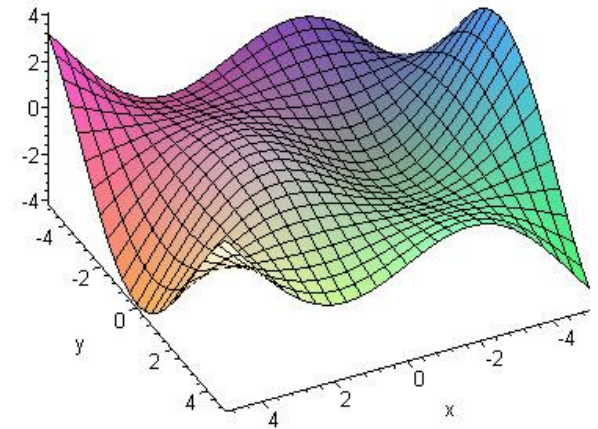
- more solution techniques for RL problems (SARSA, Q-learning)
- this time, the entire problem space does not need to be known a-priori, rather we learn to solve the problem through trial-and-error



Introduction

Gradient-based optimization

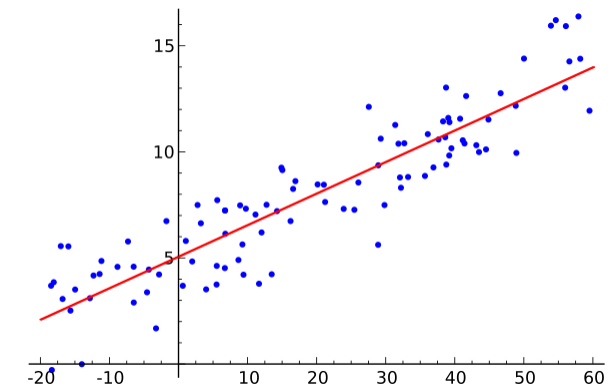
- Recap chapter for the second part of the lecture:
How to find the minimum of a differentiable function
- Some (known) concepts
(gradient, gradient descent)



https://igl.ethz.ch/teaching/tau/cg/cg2005/cg_ex6.ppt

Function approximations

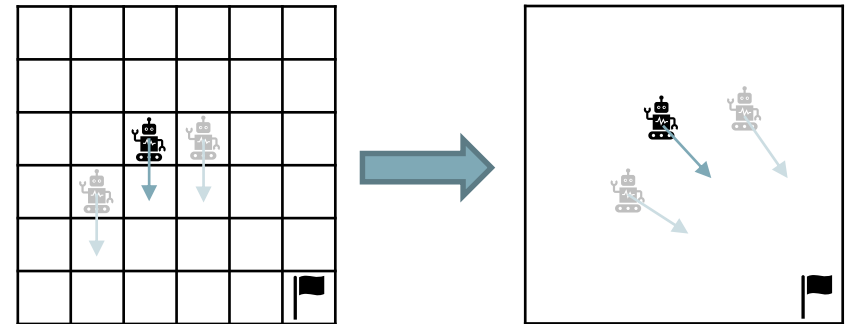
- Recap chapter for the second part of the lecture:
How function approximations can be used for RL
(e.g. neural networks)



Introduction

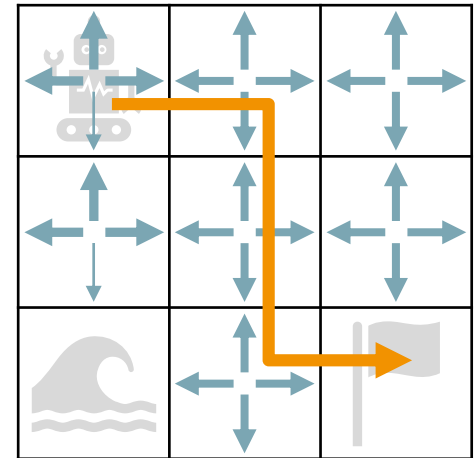
Deep RL

- Combine RL with neural networks to solve continuous environments (DQN, actor-critic)



Methods for policy-based RL

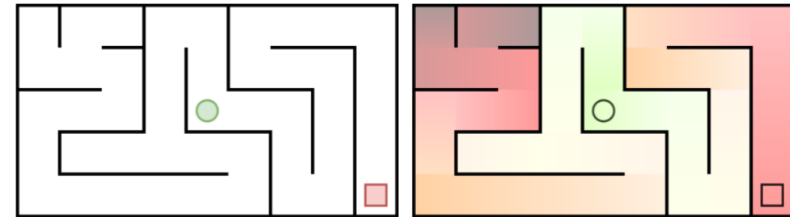
- Second large class of RL methods (REINFORCE, actor-critic)
- Constitute the most powerful methods today



Introduction

Reward shaping

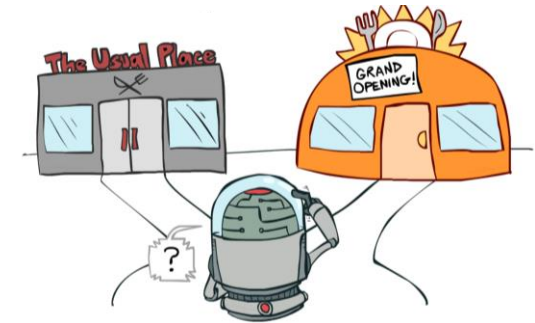
- subsidiary chapter discussing how the reward can be tweaked to improve RL performance



<https://arxiv.org/pdf/2210.09579.pdf>

Exploration and exploitation

- subsidiary chapter discussing how the agent can explore environments faster to improve RL performance

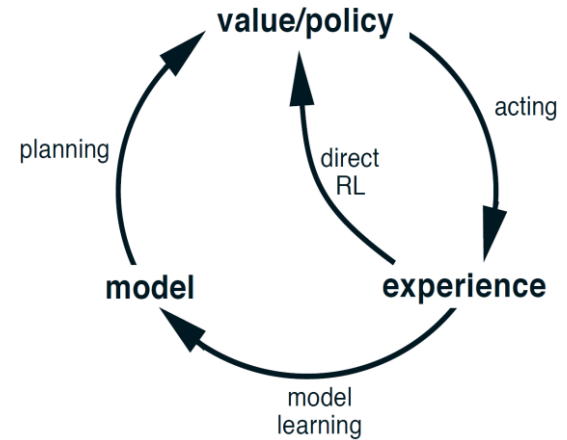


http://procaccia.info/courses/15381116/slides/781_rl_rmax.pdf

Introduction

Model-based RL

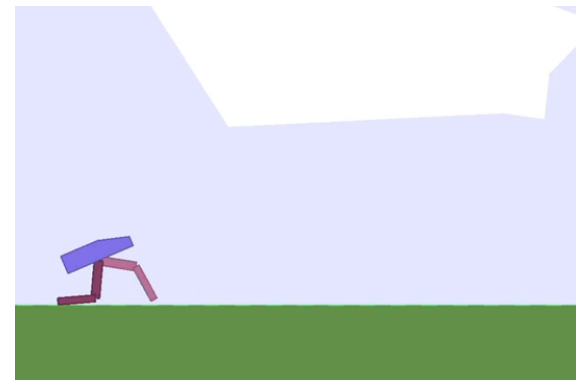
- subsidiary chapter discussing how a model of the environment can be used to improve RL performance



Sutton, Barto: Reinforcement Learning

Meta-RL

- subsidiary chapter discussing how a trained model can be applied to a different problem



<https://www.youtube.com/watch?v=D1WWWhQY9N4g>

Kahoot!

Kahoot!